STAA 551 Case Study - Home Price Regression Model

Nelson Brown and Bryce Smith

Introduction

This case study focuses on creating a regression model based on a data set of 613 transactions between the years 2014 and 2015 provided for the Seattle area in which properties of homes such as their date built, number of bedrooms, square footage of the living space, and additional features that will be described in the next section. The purpose of such a model is envisioned as an aid or input to mortgage risk analysis or to provide methodology for creating future models which may inform price negotiations between buyers and sellers today. We explore the methods of building this model up from those predictors which provide the strongest correlation with our price response variable, removing those predictors which have collinearity with others included in the model, explore interactions, and examine those models which explain the variation in the response. We perform a diagnostic and residual analysis on the best performing model to ensure it does not violate basic assumptions of the ordinary least squares (OLS) model, and utilize transformations to modify the model to fit within those assumptions at the cost of a marginal amount of explained variation. Finally, we examine the results of the final model and describe some of the uses and limitations of the model.

Our housing data frame had a number of quantitative predictors. One set are square footage of features of the house such as `sqft_living` or the amount of total space that can be lived in (sans attic, basement, or garage space), `sqft_lot` or square footage of the entire property, `sqft_basement` which is the square footage of the basement if the house has one, `sqft_above` which is living space above the basement level, `yr_built` or the year the home was built, and `yr_renovated` which is the year the house was renovated if it was indeed renovated.

A subset of the predictor variables were ordinal, discrete, or could be viewed as categorical. This included `view` (number of times a home was viewed), `grade` (assigned value by an appraiser), `condition` (another assigned value), `bedrooms`, `bathrooms`, `floors`, `waterfront`, and two predictors added in our analysis `basement` and `renovated` which were boolean values indicating if the house had those features. More information is included in the summary section of the Discrete, Ordinal, and Categorical Values.

Summary Statistics and Graphics

Table 1: Summary Statistics for Values on Seattle Housing Dataframe (2014-2015)

Statistic	Mean	St. Dev.	Min	Pctl(25)	Pctl(75)	Max
price	545,427.700	408,545.900	95,000	315,000	631,500	5,300,000
sqft_living	2,073.669	963.763	380	1,370	2,550	7,390
sqft_lot	15,967.970	46,698.890	740	5,100	10,585	871,200
sqft_above	1,793.571	873.153	380	1,130	2,313	6,530
sqft_basement	280.098	424.835	0	0	570	2,390
yr_built	1,971.210	29.939	1,900	1,951	1,998	2,015
bedrooms	3.352	0.876	1	3	4	7
bathrooms	2.092	0.805	0.500	1.500	2.500	6.000
view	0.204	0.695	0	0	0	4
grade	7.635	1.217	5	7	8	12

Summary statistics for these data points includes 613 total rows in the dataset with no missing values. This is shown in Table 1. This table gives us a summary view of our housing prices where the inter-quartile range had a midspan

between \$315k and \$630k. The minimum value of homes were approximately \$100k, and there were 154 transactions in the upper quartile of `price`. These homes were built between 1900 and 2015 being the last year our transaction data was collected.

Quantitative Values

A correlogram was generated providing the correlation matrix, density plots, and scatter plots of all relevant variables and can be seen in Figure 1. Our initial impressions or insights are given in this section. The strongest correlation between the `price` response variable and predictor variables was given by the `sqft_living` predictor at 0.723. The next highest predictor was `sqft_above`, but this had a high correlation with `sqft_living` as it is a proportion of `sqft_living`. `sqft_basement` show some amount of correlation, but this was a predictor where many homes did not have a basement so it was examined as a categorical variable which will be described in the next section. `sqft_lot` had very little correlation with `price`. The next set of predictor variables not related to area of the home was `grade` which was assigned as an ordinal value by an appraiser, `view` which was the number of times the home was viewed, and then features of the house such as `bedrooms` and `bathrooms`. Treated as ordinal or semi-ordinal these features had the next highest correlation. In the case of `bathrooms` this predictor is quasi-ordinal as a .25, .5 or .75 were really categorizations of the type of a number of bathroooms within the home. `grade` and `sqft_living` seem to have a high correlation with one another. Finally, `yr_built` and `date` of the transaction had little correlation with the `price` predictor.


Figure 1: Correlogram of Quantitative or Ordinal Predictors and Response Variable

Discrete, Ordinal, and Categorical Values

We provide Figure 2 illustrating the predictor variables that could be interpreted as factors or categorical variables. These variables are shown above. We can see a fairly clear monotonic increase in some of these ordinal values, especially the higher `grade` values. Others such as `view` and `waterfront` exhibit some of this tendency. We also added derived indicators such as `basement` and `renovated` to check if their separate distributions were largely different from one another by visual inspection and in our analysis as documented in the next section. As a note, we did conduct a separate analysis where the predictors above were treated as categorical variables and found similar or only slightly less proportion of explained variable in the adjusted R^2; however, we chose to treat our final predictors
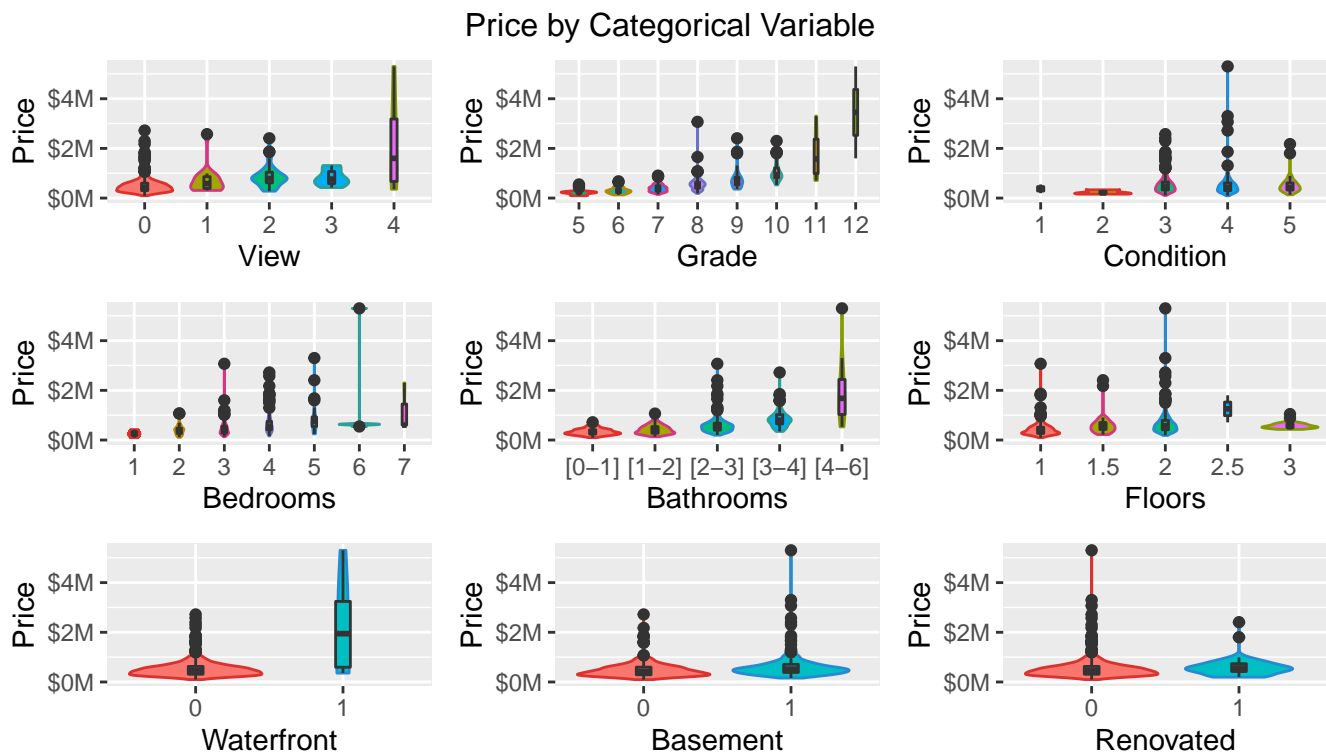
Figure 2: Violin Plots of Categorical Values Interpreted as Factors

as quantitative, ordinal values as it increased the degrees of freedom and limited predictors in the final model. This comes with the caveat that will be discussed further in the results and conclusions section that not all predictor coefficients are useful beyond the prediction of the response variable due to there not being a defined "unit-distance" between the values these ordinal predictors can take such as a half-step improvement in grade.

Analysis

Bottom-Up Model from Variable Added Last t-test

This analysis began with analyzing a correlogram of the dataset which was presented in Figure 1. Since there was a single most correlated predictor `sqft_living` with our response variable `price`, and the second highest predictor `sqft_above` was a directly proportional quantity to `sqft_living` which yielded high colinearity we decided to build a model bottom-up starting with `price` predicted by `sqft_living`. We would sequentially add predictors to the model while assessing model performance.

As seen in model (1) of Table 2, it was clear that we had strong evidence for a linear association between price and `sqft_living` and that the adjusted R^2 amount of explained variation accounted for roughly 52% of the variation in price. This would remain throughout our analysis the most significant predictor that we could ascertain contributed to the model. We then proceeded to add `grade` to the model. Adding `grade` to the model proved to be helpful, as we noticed a significant p-value for the variables added last test which was 0.000251.

Next we added bathrooms into our model and received a p-value of 4.49e-01 for the variables added last t-test so it was safe to conclude that it was not helpful to our model with `sqft_living` and `grade` already included. However, the next variable we added to our model, `view`, did prove to be helpful with a p-value of 2e-16.

The model with these three variables determine an adjusted R^2 which explained roughly 59% of the variation. We next looked at adding `bathrooms` and `bedrooms` into the model based on the correlation with `price`. `bedrooms` was chosen first as it had only a 0.429 correlation with `grade` versus the 0.740 correlation of `bathrooms` with `grade`. It was clear with the results from the summary table that bedrooms hardly explained any more variation in price, and was not considered a helpful addition in our model so we removed that predictor. `bathrooms` showed a similar result and was not included in our analysis.

We did want to try the variable `yr_built` in the model since we believed it was important to attempt to add a variable that represented the age of the house. Adding `yr_built` into our model proved to be significant with `sqft_living`, `grade` and `view` already present in the model. The results of this model are annotated as model (2) in Table 2.

At this point in our analysis we concluded to stop adding quantitative and quantitative monotonic ordinal variables, and looked at the categorical variables which could be treated as indicator or dummy variables in our model. Starting with `waterfront`, we found that this indicator was helpful in our model and increased our adjusted R^2 value to approximately 0.65.

Adding `basement` and `renovated` as indicators to our model, in the same fashion as we added `waterfront` to the model, ended up not being helpful to our model with `sqft_living`, `grade`, `view`, and `yr_built` already present. We also added these two indicators, separately, with `waterfront` in the model, and these predictors added were not helpful either. We were satisfied with the five predictors: `sqft_living`, `grade`, `view`, `yr_built`, and `waterfront` being used in our model.

Comparison of Bottom-Up Model with Top-Down Model using ANOVA tables

We wanted to ensure that our model doesn't disagree with an alternative analysis built from a full model in that we did not miss any potential predictors that may add to our model's explained variation while offering strong evidence that they were linearly associated with the predictor. To accomplish this, we first compared a model with nearly every predictor with the exception of those predictors that could be eliminated due to their intrinsic nature or prior inspection. For instance, the id field or date of transaction was not considered in the full model. An ANOVA was conducted between these models with an F-test statistic value of 3.969 and associated p-value 6.64e-04 indicated that their may be some predictors which may be linearly associated with the response variable.

We then removed the predictors from the nearly full model by examining those with the weakest p-value evidence, and ended up with a reduced model which discarded `floors`, `sqft_above`, and `sqft_lot`. The reduced model when compared with the full model through an ANOVA table had a F-test statistic of 2.002 and associated p-value of 0.093 which gave fair evidence that the removed predictors were not linearly associated with `price`. Comparing the adjusted R^2 value of this reduced model of 0.661 with the model we arrived at built from a bottom-up analysis using variable last added t-tests from initial predictors and it's adjusted R^2 of 0.621 indicates that the additional explained variation is minimal compared with our model with fewer predictors. A weakness of this process is that the ANOVA analysis is order dependent on the predictors, but conducting it gives some amount of confidence in the model.

Exploring Interactions

Proceeding with model (2) in Table 2, we needed to assess the correlations between our predictors. When doing this, it was clear that we had evidence of colinearity between `sqft_living` and `grade`. With this knowledge we sought to assess whether an interaction between these two variables would be helpful to our model. As seen in model (3) in Table 2, adding this interaction term proved to be helpful in our model and increased the variation in `price` explained by the model to roughly 71% as expressed in the adjusted R^2. At this point, we moved on to the diagnostic plots and residuals analysis. Models (4) and (5) in Table 2 will be explained in the following sections.

Table 2: Iteratively Built Model Comparison

	Dependent variable:				
	price			log(price)	
	(1)	(2)	(3)	(4)	(5)
sqft_living	306.516***	210.259***	−338.813***		
	(11.846)	(17.565)	(52.678)		
log(sqft_living)				0.185	0.426***
				(0.147)	(0.046)
grade		103,280.400***	−22,648.160	−0.025	0.236***
		(14,665.420)	(18,045.770)	(0.153)	(0.018)
view		128,209.200***	63,479.420***	0.109***	0.115***
		(15,686.520)	(16,038.800)	(0.023)	(0.020)
yr_built		−2,834.238***	−2,411.571***	−0.004***	−0.005***
		(406.927)	(361.830)	(0.001)	(0.001)
waterfront			668,962.300***	0.026	
			(110,486.600)	(0.157)	
sqft_living:grade			61.812***		
			(5.796)		
log(sqft_living):grade				0.033*	
				(0.019)	
Constant	−90,183.990***	4,881,654.000***	5,119,662.000***	18.658***	16.976***
	(27,085.020)	(764,312.600)	(673,590.900)	(1.387)	(0.992)
Observations	613	613	613	613	613
R^2	0.523	0.624	0.709	0.648	0.646
Adjusted R^2	0.522	0.621	0.706	0.644	0.643
Residual Std. Error	282,443.100 (df = 611)	251,492.200 (df = 608)	221,542.500 (df = 606)	0.319 (df = 606)	0.320 (df = 608)
F Statistic	669.475*** (df = 1; 611)	251.761*** (df = 4; 608)	245.871*** (df = 6; 606)	185.617*** (df = 6; 606)	277.155*** (df = 4; 608)

Note: *p<0.1; **p<0.05; ***p<0.01

Diagnostic Plots and Residuals Analysis

Satisfied with the model derived up to this point, we proceeded to the diagnostic plots shown in Figure 3 to check the validity of our assumptions. After review of these plots, it was clear we had problems addressing our assumptions of linearity, mean zero for the errors and constant variance of the errors. This can be shown with the fanning out of the points in the residuals vs fitted plot, and the lack of linearity in the normal QQ plot. We can also see that we have observations with very large Cook's Distance, which may be negatively affecting our model. Checking the residuals versus the predicted values confirmed what we saw in the previous three plots and we concluded that our assumptions were violated.

Figure 3: Diagnostic Plots

To address the violations of our assumptions, we tried transforming some of the variables used in the model. We could see from the correlogram that both our response, `price`, and predictor, `sqft_living`, had heavily right skewed distributions. We decided to apply a logarithmic transformation after observing the shift in distribution indicating that the logarithmic transformation would be justified and this is shown in Figure 4. We reapplied these into our model.

Doing so we saw our interaction between the now logarithmic transformation of `sqft_living` and `grade` was no longer helpful to our model. We also saw that the indicator, `waterfront`, no longer displayed strong evidence of linear association. This can be seen in model (4) of Table 2. In an effort to satisfy our assumptions of the model, it was decided that removing both of these predictors from the model was necessary. This is shown in our final model (5) of Table 2. As a quick check we looked into the variable inflation factor (VIF) values for the predictors and determined that we weren't worried about collinearity between the predictors shown in Table 3. We rexamined the diagnostic plots which look to satisfy our base assumptions of the analysis shown in Figure 5.

Figure 4: Transformations of Skewed Distributions

Table 3: VIF for Final Price Model

log(sqft_living)	grade	view	yr_built
2.543	2.874	1.112	1.438

Lastly, in an effort to address the values that showed large Cook's Distance in the diagnostic plots, we compared which leverage points in our model also had large residuals. We found observations that satisfied both these criteria, and concluded that these points needed to be addressed with a subject matter expert if one were available.

Results and Conclusions

Our final fitted model is given by the following regression equation:

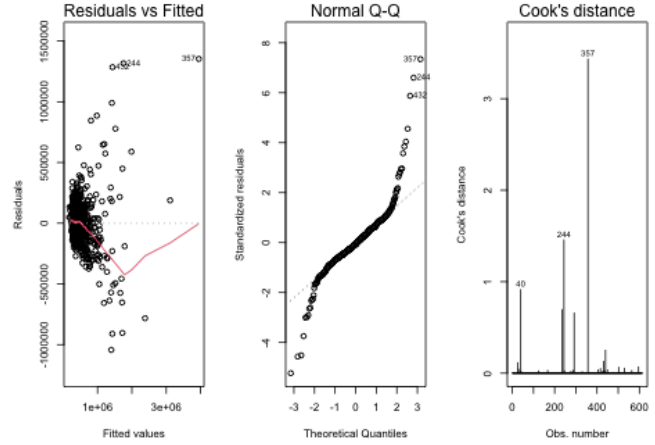$$\log(price) = 16.98 + 0.43\log(sqftliving) + 0.24grade + 0.12view - 0.0045yr\_built$$
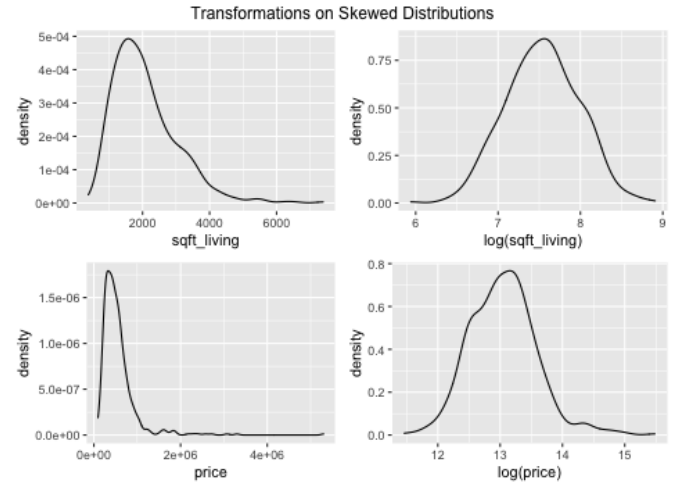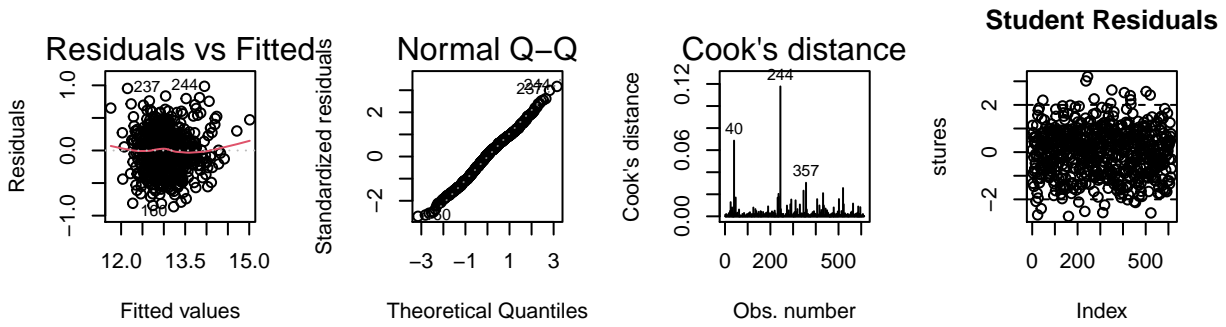
Figure 5: Diagnostic Plots of Final Model

Table 4: High Student Residual and Leverage Values

	price	sqft_living	view	grade	yr_built
40	357,000	2,460	4	7	1,955
47	550,000	1,660	0	5	1,933
244	3,070,000	3,930	4	8	1,957
346	245,000	380	0	5	1,963
432	2,720,000	3,990	0	11	1,989
520	675,000	930	2	6	1,951

Full details are available as model (5) in Table 2. An ANOVA table for this is provided at the end of this section. This model is likely suited for comparative analysis of home prices and their attributes between 2014 and 2015 in the Seattle area, where one may extend that time period a few months (up to a year) before or after but likely not outside of that spatial extent. It would likely be useful in creating a representative prediction used for mortgage risk analysis or as a method to create and inform a negotiated sale price model with today's data. Some home buyers or sellers might use this regression equation to examine actionable or significant features of their homes with some caveats.

The most actionable information from the coefficients of this equation is the logarithm of `sqft_living`. Holding all else constant, a 10 percent increase in the living square footage of a home would yield a 4.26% increase in the average price of the home. `yr_built` has a negative coefficient which can be interpreted as many older homes in the Seattle area were generally in a condition that drove higher sales. `grade` and `view` should be used for analysis carefully. Holding all else constant, a one level increase in `grade` or `view` will result in an increase in average price of 26.67% and 12.23%, respectively. But likely there are attributes not captured in this model which would make that information actionable. `grade` is likely a very subjective measure based on hidden variables (cost of cabinet remake, or modern/ traditional style) and the number of times a home is viewed is likely dependent on those same hidden variables. This would likely need it's own modelling to ascertain how homeowners or buyers could make that information actionable, and in the context of this model it is likely providing value in its predictive power after the value of those predictor variables are known.

Finally, an examination of the ANOVA table below indicates that there is strong evidence that all of these predictors are linearly associated with our response variable `price`. We believe that this model is useful when used within the limitations and purposes noted.

```
## Response: log(price)
##                  Df Sum Sq Mean Sq F value    Pr(>F)
## log(sqft_living)  1 86.819  86.819 849.926 < 2.2e-16 ***
## grade             1 11.943  11.943 116.921 < 2.2e-16 ***
## view              1  6.587   6.587  64.487 5.049e-15 ***
## yr_built          1  7.895   7.895  77.287 < 2.2e-16 ***
## Residuals       608 62.106   0.102
```

Appendix: Reference and All Code for This Report

Hlavac, Marek (2018). stargazer: Well-Formatted Regression and Summary Statistics Tables. R package version 5.2.1. https://CRAN.R-project.org/package=stargazer

```r
knitr::opts_chunk$set(echo = FALSE, warning = FALSE, dev='pdf')
library(car)
library(ggplot2)
library(grid)
library(gridExtra)
library(ggthemes)
library(lubridate)
library(GGally)
library(scales)
library(dplyr)
library(stargazer) # Used for latex tables to summarize the data and models

# SECTION: HEADER AND INTRODUCTION, INITIAL DATA READ

# Helper function for formatting
sci.notation <- function(value) {
  formatC(value, format = "e", digits = 2)
}

# Read the Data
df <- read.csv('Seattle.csv', strip.white = TRUE, stringsAsFactors = FALSE)
# Clean the Data
df$date <- ymd(substr(df$date,1,nchar(df$date) - 7)) # Convert string to date object
df$was_renovated <- as.factor(with(df,
    ifelse(yr_renovated > 0,
          1,
          0
        )
 ))
df$has_basement <- as.factor(with(df,
    ifelse(sqft_basement > 0,
          1,
          0
        )
 ))

# SECTION: SUMMARY STATISTICS AND GRAPHCIS

# Define our quantitative values of interest
df.quant <- df %>% dplyr::select(price,
                       sqft_living,
                       sqft_lot,
                       sqft_above,
                       sqft_basement,
                       yr_built,
                       bedrooms,
                       bathrooms,
                       view,
                       grade,
                       date)

# Summarize initial dataframe
```

```r
stargazer(df.quant,
          header=FALSE,
          omit.summary.stat=c('N'),
          title="Summary Statistics for Values on Seattle Housing Dataframe (2014-2015)",
          label="summary-stats")


# SECTION: QUANTITATIVE VALUES

# Quantitative Data Pairs
ggpairs(df.quant, progress=FALSE) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1))


# SECTION: DISCRETE AND CATEGORICAL VALUES

# Binning bathrooms
tags <- c("[0-1]", "[1-2]","[2-3]", "[3-4]", "[4-6]")
v <- df %>% dplyr::select(price, bathrooms)
vgroup <- as_tibble(v) %>% mutate(tag = case_when(
  bathrooms == 0.25|bathrooms == 0.5|bathrooms == 0.75|bathrooms == 1.00 ~ tags[1],
  bathrooms == 1.25|bathrooms == 1.5|bathrooms == 1.75|bathrooms == 2.00 ~ tags[2],
  bathrooms == 2.25|bathrooms == 2.5|bathrooms == 2.75|bathrooms == 3.00 ~ tags[3],
  bathrooms == 3.25|bathrooms == 3.5|bathrooms == 3.75|bathrooms == 4.00 ~ tags[4],
  bathrooms > 4.00 & bathrooms <= 6.00 ~ tags[5],
))
df$bathroom_groups <- as.factor(vgroup$tag)


plot_price_by_cat <- function(df, cat_var, cat_var_name) {
  ggplot(df, aes(x=cat_var, y=price, fill=cat_var)) +
    scale_colour_solarized("red") +
    geom_violin(aes(color=cat_var)) +
    geom_boxplot(width=0.1) +
    xlab(cat_var_name) +
    ylab("Price") +
    scale_y_continuous(labels = scales::dollar_format(scale = .000001, suffix="M")) +
    theme(legend.position="none")
}

q <- plot_price_by_cat(df, as.factor(df$view), "View")
r <- plot_price_by_cat(df, as.factor(df$grade), "Grade")
x <- plot_price_by_cat(df, as.factor(df$condition), "Condition")
b <- plot_price_by_cat(df, as.factor(df$bedrooms), "Bedrooms")
bath_plot <- plot_price_by_cat(df, df$bathroom_groups, "Bathrooms")
w <- plot_price_by_cat(df, as.factor(df$floors), "Floors")
t <- plot_price_by_cat(df, as.factor(df$waterfront), "Waterfront")
y <- plot_price_by_cat(df, as.factor(df$has_basement), "Basement")
reno <- plot_price_by_cat(df, df$was_renovated, "Renovated")

grid.arrange(grobs=list(q, r, x,
                        b, bath_plot, w,
                        t, y, reno),
             ncol=3,
             top="Price by Categorical Variable")

# SECTION BOTTOM-UP MODEL FROM VARIABLE ADDED LAST t-TEST
```

```r
lm.price.1 <- lm(formula = price ~ sqft_living, data = df)
lm.price.2 <- lm(formula = price ~ sqft_living + grade, data=df)
lm.price.3.not.included <- lm(formula = price ~ sqft_living + grade + bathrooms,
                              data = df)
lm.3.pvalue <- sci.notation(summary(lm.price.3.not.included)$coefficients[4,4])
lm.price.4 <- lm(formula = price ~ sqft_living + grade + view,
         data = df)

lm.price.5 <- lm(formula = price ~
                 sqft_living +
                 grade +
                 view +
                 yr_built,
         data = df)

# Adding categorical predictors

lm.price.5b <- lm(formula = price ~
                 sqft_living +
                 grade +
                 view +
                 yr_built +
                 waterfront,
               data = df)

# SECTION: COMPARISON OF BOTTOM-UP MODEL WITH TOP-DOWN MODEL USING ANOVA

# Top-down model building using ANOVA analysis between a nearly
# full model and our bottom-up analysis, a nearly full model
# and a reduced model, and finally comparative metrics between
# the reduced model and the model built from a bottom-up analysis.


lm.nearfull <- lm(price ~
                 bedrooms +
                 bathrooms +
                 sqft_living +
                 sqft_lot +
                 floors +
                 waterfront +
                 view +
                 condition +
                 grade +
                 sqft_above +
                 yr_built , data=df)
anova.1 <- anova(lm.nearfull, lm.price.5b)
anova.1.p.value <-formatC(anova.1$`Pr(>F)`[2], format = "e", digits = 2)
lm.reduced <- lm(price ~ bedrooms +
                 bathrooms +
                 sqft_living +
                 waterfront +
                 view +
                 grade +
                 yr_built, data=df)
anova.2 <- anova(lm.nearfull, lm.reduced)
anova.2.p.value <-round(anova.2$`Pr(>F)`[2], 3)
```

```r
# collect adjusted R^2
nearfull.r.adjusted <- summary(lm.nearfull)
reduced.r.adjusted <-round(summary(lm.nearfull)$adj.r.squared,3)
lm5.r.adjusted <-round(summary(lm.price.5)$adj.r.squared,3)


# SECTION: EXPLORING INTERACTIONS

lm.price.6 <- lm(formula = price ~
                 sqft_living +
                 grade +
                 view +
                 yr_built +
                 waterfront +
                 sqft_living*grade,
                 data = df)

# Corrections and transformations applied based on diagnostic plots
# (performed before diagnostic plots as to provide input to table 4
#  used in the report which we didn't want separated too far from
#  the bottom-up model building section of analysis)

lm.price.7 <- lm(formula = log(price) ~
                 log(sqft_living) +
                 grade +
                 view +
                 yr_built +
                 waterfront +
                 log(sqft_living)*grade,
                 data = df)
# Final model
lm.price.8 <- lm(formula =
                 log(price) ~
                 log(sqft_living) +
                 grade +
                 view +
                 yr_built,
                 data = df)
lm.price.final <- lm.price.8
stargazer(lm.price.1,
          lm.price.5,
          lm.price.6,
          lm.price.7,
          lm.price.8,
          header=FALSE,
          align=T,
          label="ModelComparison",
          title="Iteratively Built Model Comparison")

# SECTION: DIAGNOSTIC PLOTS AND RESIDUAL ANALYSIS

# In order to save space, we had to flow some text around these
# diagnostic figures, but I think it worked well to keep the
# figures near their explanation.

png("diagnostics_prior_to_transform.png", height = 350)
```

```r
# 2. Create the plot
par(mfrow = c(1,3))
plot(lm.price.6,c(1,2,4))
# 3. Close the file
dev.off()
# Box-cox examining most likely model predictor exponent
png("boxcox.png", height = 350)
boxCox(lm.price.6, main="Box-Cox Plot of model")
dev.off()

png("transformations.png", height = 350)
par(mfrow = c(1,2))
p <- ggplot(df, aes(x=sqft_living)) + geom_density()
q <- ggplot(df, aes(x=log(sqft_living))) + geom_density()
r <- ggplot(df, aes(x=price)) + geom_density()
s <- ggplot(df, aes(x=log(price))) + geom_density()
grid.arrange(grobs=list(p, q,
                        r, s),
             ncol=2,
             top="Transformations on Skewed Distributions")
dev.off()

# VIF Table
out <- capture.output(stargazer(vif(lm.price.final),
         title="VIF for Final Price Model",
         header=FALSE,
         align=FALSE))

out <- sub(" \\\\centering", "", out)
cat(out)
# Student Residuals and Hat Values Computed
par(mfrow = c(1,5))
plot(lm.price.final,c(1,2,4))
stures <- rstudent(lm.price.final)
plot(stures, main="Student Residuals")
abline(h=2, lty=2)
abline(h=-2, lty=2)
# Calculate intersection of student residuals with high leverage
hv <- as.matrix(hatvalues(lm.price.final))
mn <- mean(hatvalues(lm.price.final))
stures_vals <- which(abs(stures) > 2)
lev_vals <- which(hv > 2*mn)
index_values <- intersect(lev_vals,stures_vals)
df.new <- df %>% dplyr::select(price,
                        sqft_living,
                        view,
                        grade,
                        yr_built)
stargazer(df.new[index_values,],
         header=FALSE,
         title="High Student Residual and Leverage Values",
         summary=FALSE)

# Final ANOVA table presented in results and conclusions

out <- capture.output(anova(lm.price.final))
```

```r
cat(paste(out[3:9],collapse='\n'))
```